

Artificial intelligence


1. What exactly is AI?

Carl McBride Ellis

 carl.mcbride@u-tad.com


Recommended bibliography

- François Chollet “*Deep Learning with Python*” 3rd Edition, Manning (2025)
- Free online version
- GitHub repo



Recommended bibliography

- Aurélien Géron "*Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow*", 3rd Edition O'Reilly Media (2022) (especially Part II)
- GitHub repo




Recommended bibliography

- A Comprehensive Overview and Comparative Analysis on Deep Learning Models: CNN, RNN, LSTM, GRU

Before 1956...

- ~~cybernetics~~
- ~~automata theory~~
- ~~information processing~~
- ~~thinking machines~~
- ~~statistical computing~~
- ~~computational statistics~~
- ~~predictive analytics~~

“artificial intelligence”



Source: John McCarthy, Marvin L. Minsky, Nathaniel Rochester, Claude E. Shannon "A *Proposal for the Dartmouth Summer Research Project on Artificial Intelligence*", August 31 (1955)

1956 Dartmouth Conference: The Founding Fathers of AI




John McCarthy



Marvin Minsky



Claude Shannon



Ray Solomonoff




Alan Newell



Herbert Simon



Arthur Samuel



Oliver Selfridge




Nathaniel Rochester



Trenchard More

...and now we have:



Lynx A.I. Limited Edition deodorant. Fragrance powered by Artificial Intelligence. Get the code to smell iconic.

ALL NEW LYNX A.I.
LIMITED EDITION

Our first limited edition. Lynx AI deodorant, powered by Artificial Intelligence. 46 Terabytes of data, 6,000 ingredients, 3.5 million possible combinations, 1 fragrance. Do you want the code to smelling iconic?

BUY IT NOW



Coca-Cola CREATIONS Y3000

LIMITED EDITION FLAVOR TASTE THE FUTURE

Co-Created with AI

Y3000 ZERO SUGAR TANPA GULA

© 2023 The Coca-Cola Company



INTELIGENCIA ARTIFICIAL en el DISEÑO DEL VINO

Cata "AIWine"

SÁBADO 18 / NOVIEMBRE / 2023

HORARIO: De 11:30 a 14 h.
LUGAR: Oficina de Turismo de la Acera de Recoletos (Valladolid)
Plazas limitadas. Asistencia gratuita previa inscripción.

UEMC UNIVERSIDAD
COLABORADOR:
Ayuntamiento de Valladolid
www.valladolid.es

SEMANA DE LA CIENCIA EN CASTILLA Y LEÓN
MOVIMIENTO DE 2023

6. Artificial intelligence



When the function is activated, the temperature of the refrigerator is automatically switched to

5°C, the temperature of freezer is automatically switched to -18°C

however

4. EXAMPLES OF SYSTEMS WITH MENTAL QUALITIES

Let us consider some examples of machines and programs to which we may ascribe belief and goal structures.

104 FORMALIZING COMMON SENSE

4.1. *Thermostats*

Ascribing beliefs to simple thermostats is unnecessary for the study of thermostats, because their operation can be well understood without it. However, their very simplicity makes it clearer what is involved in the ascription, and we maintain (partly as a provocation to those who regard attribution of beliefs to machines as mere intellectual sloppiness) that **the ascription is legitimate.**⁷

Source: *Ascribing Mental Qualities to Machines*, John McCarthy (1979)

3 Types of Artificial Intelligence

Artificial Narrow Intelligence (ANI)



Stage-1

Machine Learning

- ▶ Specialises in one area and solves one problem



Artificial General Intelligence (AGI)



Stage-2

Machine Intelligence

- ▶ Refers to a computer that is as smart as a human across the board

Artificial Super Intelligence (ASI)



Stage-3

Machine Consciousness

- ▶ An intellect that is much smarter than the best human brains in practically every field

François Chollet “On the Measure of Intelligence”



Elon Musk

@elonmusk

Hope we're not just the biological boot loader for digital superintelligence. Unfortunately, that is increasingly probable

9:18 PM · Aug 3, 2014

The Chollet Venn diagram (awful!)




Figure 1.1 Artificial intelligence, machine learning, and deep learning

The EU definition of AI:

Part of [Chapter I: General Provisions](#)

Article 3: Definitions

Date of entry into force:

2 February 2025

According to:

[Article 113\(a\)](#)

Inherited from:

[Chapter I](#)


See here for a [full implementation timeline](#).

SUMMARY +

For the purposes of this Regulation, the following definitions apply:

- (1) 'AI system' means a machine-based system that is designed to operate with varying levels of autonomy and that may exhibit adaptiveness after deployment, and that, for explicit or implicit objectives, infers, from the input it receives, how to generate outputs such as predictions, content, recommendations, or decisions that can influence physical or virtual environments; Related: [Recital 12](#)

I personally would (humbly) suggest:



TL;DR:

Machine learning is the (objective) technique that is used

Whether it is AI or not is based on the (subjective) use-case

- tabular data = ML
- image data = Artificial (narrow) Intelligence (AI)
- text data = Artificial (narrow) Intelligence (AI)
- image + text = Artificial General Intelligence (AGI?)
- ¿what is **Artificial SuperIntelligence?** (ASI)

Turing test (1950): The ‘Imitation Game’

“The new form of the problem can be described in terms of a game which we call the ‘imitation game’. It is played with three people, a man (A), a woman (B), and an interrogator (C) who may be of either sex. The interrogator stays in a room apart from the other two. The object of the game for the interrogator is to determine which of the other two is the man and which is the woman.”

Alan M. Turing “*Computing Machinery and Intelligence*”, Mind vol LIX pp. 433-460 (1950)

- Loebner Prize 1991-2019

\$25,000 prize for the first bot that can pass a text-only Turing Test, and \$2,000-\$3,000 for the most human-seeming of all contestants

The ‘Lady Lovelace objection’ (1843):

“The Analytical Engine has no pretensions whatever to originate anything. It can do whatever we know how to order it to perform”.

Bernardo Gonçalves [*“Lady Lovelace’s Objection: The Turing–Hartree Disputes Over the Meaning of Digital Computers, 1946–1951”*](#), IEEE Annals of the History of Computing, vol. 46 pp. 6-18 (2024)

Jacquard machine punch cards for the fabric design (1750 - 1800)




The Babbage Difference engine (c. 1822)




(See: "[Charles Babbage's Difference Engines and the Science Museum](#)")

and today:




different, ...but the same.

The Chinese room



(John R. Searle "Minds, brains, and programs" (1980))

The five tribes of machine learning



taken from Pedro Domingos "*The Master Algorithm*" (2015)

- symbolists - logic-based artificial intelligence
- evolutionaries - genetic programming
- Bayesians - probabilistic graphical models
- analogizers - classical machine learning
- connectionists - artificial neural networks

The Bayesian **prior probability distribution** can be ill-posed: perfect ignorance demands a uniform prior. However:

“...if we are considering an unknown square, we could say we are uniformly ignorant of the square’s side length (S) or its area (A) - but not both. Since the two measurements are related by the function $A = S^2$, they cannot both be uniformly distributed, but either could reasonably serve as a description of the size of the square”.

analogizers → classical machine learning (tabular data)


- regression / classification
- decision trees
- Random Forest
- Gradient Boosted Decision Trees (GBDT)

used for point predictions and confidence intervals

connectionists → artificial neural networks (ANN) / “deep learning”

- computer vision (CV) with CNN
 - image classification
 - image segmentation
 - image-to-text (OCR: optical character recognition)
- natural language processing (NLP) with RNN
 - sentiment analysis
 - language translation
 - speech-to-text (STT)
- time series (LSTM)

example of image segmentation



(segmentation is classification with a 'void' class)

The 'Gen-AI IPO' of 2022

[Stable Diffusion](#) by Stability AI - August 22, 2022

[DALL-E 2](#) by OpenAI - September 28, 2022

[ChatGPT](#) by OpenAI - November 30, 2022




(paper: "[Tutorial on Diffusion Models for Imaging and Vision](#)", 6 September 2024)

examples of generative AI (Gen-AI) applications


- text-to-text: prompts, chatbots ([ChatGPT](#)), translation ([DeepL](#))
- text-to-speech: [ElevenLabs](#)
- text-to-image: [Dall-E 3](#), [Midjourney](#), [Stable Diffusion](#), [Flux.1](#)
- image-to-text: [image captioning](#)

Dominant Generative AI Use Cases



Source: "2024: The State of Generative AI in the Enterprise" (November 20, 2024)

Gen-AI landscape




- transformer models (2017)
paper: “Attention Is All You Need” by Google
- foundation models (2021)
paper: “A Comprehensive Survey on Pretrained Foundation Models:
A History from BERT to ChatGPT”

Large language models (LLM)

- BERT (Bidirectional Encoder Representations from Transformers)
Google (2018)
- GPT-1 (Generative Pre-trained Transformer 1)
OpenAI (2018)

LLM timeline:



Reviews:

- “A Survey of Large Language Models” (November 2023)
- “Large Language Models: A Survey” (February 2024)
- “A Comprehensive Overview of Large Language Models” (April 2024)

The ELIZA effect - anthropomorphism

- ELIZA is a chatbot from 1966

paper "ELIZA—a computer program for the study of natural language communication between man and machine" (1966)

Sample ELIZA output:

ELIZA: How do you do. Please tell me your problem.

YOU: I am worried about AI

ELIZA: Is it because you are worried about ai that you came to me ?

YOU: Yes, that is why

ELIZA: You seem to be quite positive.

YOU: I like to try to be positive

ELIZA: That's quite interesting.

YOU: I am glad you think so

ELIZA: How have I helped you to be glad ?

YOU: By listening

ELIZA: That is interesting. Please continue.

Try ELIZA yourself!

(see also the [ELIZA source code in BASIC](#))

pareidolia:



LLMs are as persuasive as humans; linguistic fluency is mistaken for epistemic reliability:

“...LLMs’ persuasive abilities stem from their strategic use of grammatical and lexical complexity as well as their deployment of moral language, underscoring the models’ capability to deeply engage users and invoke significant cognitive and moral processing.”

“Large Language Models are as persuasive as humans, but how? About the cognitive effort and moral-emotional language of LLM arguments.”

Microsoft researchers think GPT-4 shows AGI

paper "["Sparks of Artificial General Intelligence: Early experiments with GPT-4"](#)" (March 2023)


However, LLM's lack a [world model](#)

LLM's == stochastic parrot:

	A parrot	Machine learning algorithm
Learns random phrases	✓	✓
Doesn't understand shit about what it learns	✓	✓
Occasionally speaks nonsense	✓	✓
Is a cute birdie parrot	✓	✗

(paper: "On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?", 2021)

“Everything an LLM outputs is a hallucination. It’s just that some of those hallucinations are true.”



See also “[Why Language Models Hallucinate](#)” (Open AI, September 2025)

The AI singularity

“...an upgradable intelligent agent could eventually enter a positive feedback loop of self-improvement cycles, each successive; and more intelligent generation appearing more and more rapidly, causing a rapid increase (“explosion”) in intelligence which would ultimately result in a powerful superintelligence, qualitatively far surpassing all human intelligence”

(Source: Wikipedia)

Dead Internet theory (c. 2021)

“The dead Internet theory is an online conspiracy theory that asserts that the Internet now consists mainly of bot activity and automatically generated content manipulated by algorithmic curation to intentionally manipulate the population and minimize organic human activity.”

(Source: Wikipedia)

September 5th 2025:



Futurism

OH REALLY? | SEP 5, 8:00 AM EDT by FRANK LANDYMORE

Sam Altman Says He's Suddenly Worried Dead Internet Theory Is Coming True

"We're all trying to find the guy who did this!"

paper: “*A Shocking Amount of the Web is Machine Translated*”

paper: “*AI models collapse when trained on recursively generated data*”

(paper “Datasets for Large Language Models: A Comprehensive Survey” February 2024)

AI winters

THE RISE OF AI




Figure 1: An AI timeline; Source: Lavenda, D./Marsden, P.

source dhl via @mikequindazzi

- State of Data Science survey Anaconda (2023)
- State of AI Report 2023
- Kaggle AI Report 2023



commentary in snack-sized mouthfuls

"I want AI to do my laundry and dishes so that I can do art and writing, not for AI to do my art and writing so that I can do my laundry and dishes."


Author and videogame enthusiast **Joanna Maciejewska** nails it
(although bathroom cleaning goes ahead of laundry and dishes)

"I'm sure I deserve a lot of..."


Course overview

- Introduction to python, Jupyter Notebooks, Google Colab and Kaggle
- **tabular data:** ANN
 - linear regression
 - logistic regression and classification
 - artificial neural networks using Keras
 - single neuron: linear regression
 - single neuron: logistic regression
- **image data:** CNN
 - history: ImageNet competition and AlexNet
 - practicals: MNIST / fashion MNIST / CIFAR-10
 - pre-trained Keras models: VGG-16, ResNet and EfficientNet
 - **Kaggle project: binary image classification ¿cat or dog? (20%)**
- **text data:** RNN
 - cleaning with REGEX, scraping with BeautifulSoup
 - tokenization and vectorization (BoW, TF-IDF,...)
 - recurrent neural networks: LSTM with Keras
 - **Kaggle project: binary text classification SPAM/NOT_SPAM (20%)**
- unsupervised deep learning
 - autoencoders
 - GAN
- reinforcement learning
 - multi-armed bandit
- Generative-AI
 - the Gen-AI ecosystem
 - Agentic AI
- Ethics, AI and society, GDPR and the EU AI Act
- **Final individual project (20%)**

Machine learning: tabular data



(some) machine learning libraries:



for (classical) machine learning



for time series



but there are many, many more...

ML, AI and Data landscape 2024

