


# Linear regression

Carl McBride Ellis

✉ carl.mcbride@u-tad.com

## Regresión lineal: Ejemplo en 1D ( “regresión lineal univariante”):



## The linear regression model:

$$\hat{y} = \beta_1 x + \beta_0 + \epsilon \text{ we assume that } \mathbb{E}(\epsilon) = 0, \text{ p.ej. } \sim \mathcal{N}(0, \sigma^2)$$

For each '*feature*'  $x$  there is an adjustable parameter,  $\beta_1$ , which corresponds to its slope or gradient (often called the **weight**) as well as a intersection term  $\beta_0$  (often called the **bias**).

This is a 'parametric' model; here we assume the relation between the dependent and the independent variables is linear, and the noise is from a (**Gaussian**) parametric distribution.

Our data, especially when working with tabular data, is usually stored in a DataFrame (`df`):

	feature_1	feature_2	feature_3	feature_4	target
0	-0.285624	-1.109467	-0.704222	1.311941	1.011175
1	0.115344	-0.582376	0.850097	-0.078389	-1.017000
2	0.168309	0.781371	0.564461	1.798512	-0.339078
3	-0.364210	0.427111	1.909429	1.822293	-0.702626
4	-0.054081	2.215833	-0.384593	0.634965	-1.837298
5	-0.104199	0.828819	0.553767	-0.325654	0.885174
6	-0.143291	-1.508430	-0.183994	0.063684	0.234775
7	-1.076164	-0.233883	0.404790	-1.287040	-1.522206
8	0.209559	-0.778077	-0.207529	0.560157	-0.074932
9	0.222729	-0.549980	-0.459486	1.581879	0.921162

(In deep learning is it often stored in `numpy arrays` or in `tensors`).

## Solver: Ordinary least squares (OLS)

$$y = \beta_1 x + \beta_0$$

$$\beta_1 = \frac{\sum((x_i - \bar{x})(y_i - \bar{y}))}{\sum((x_i - \bar{x})^2)}$$

$$\beta_0 = \bar{y} - \beta_1 \bar{x}$$

(The technique of least squares was invented by [Carl Friedrich Gauss](#) and [Adrien-Marie Legendre](#) around 1802-1809)

# Practical session

Practical session: linear regression with python

[`notebook\_ML\_OLS\_numpy.ipynb`](#)

Objective: convert the equations for ordinary least squares into python code, and visually inspect the result

Our DataFrame is basically an **augmented coefficient matrix**:

		total
3 	3 	21€
5 	2 	20€
8 	6 	46€

(Note: this example is overdetermined)

Say we were paying by card, and there was also a fixed transaction fee of 2€

apple	grapes	coin	total
3 apple	3 grapes	1 coin	23€
5 apple	2 grapes	1 coin	22€
8 apple	6 grapes	1 coin	48€

This constant term would be our  $\beta_0$

We can treat our DataFrame as if it were a system of linear equations where the values of  $x_{mn}$  represent the coefficients:

$$\begin{cases} \beta_1 \cdot x_{0,1} + \beta_0 \cdot 1 = y_0 \\ \beta_1 \cdot x_{1,1} + \beta_0 \cdot 1 = y_1 \\ \beta_1 \cdot x_{2,1} + \beta_0 \cdot 1 = y_2 \end{cases}$$

where we wish to find the  $\beta_0$  and  $\beta_1$  associated with each column.

(The constant column can be created using `statsmodels add_constant`)

Written as matrices:

$$\mathbf{X} = \begin{bmatrix} 1 & x_0 \\ 1 & x_1 \\ \vdots & \vdots \\ 1 & x_m \end{bmatrix}, \quad \boldsymbol{\beta} = \begin{bmatrix} \beta_0 \\ \beta_1 \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_0 \\ y_1 \\ \vdots \\ y_m \end{bmatrix}$$

( $\mathbf{X}$  is known as the '*design matrix*')

## In more dimensions

Multiple linear regression (now with  $n$  features):

$$\mathbf{X} = \begin{bmatrix} 1 & x_{01} & x_{02} & \cdots & x_{0n} \\ 1 & x_{11} & x_{12} & \cdots & x_{1n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_{m1} & x_{m2} & \cdots & x_{mn} \end{bmatrix}, \quad \beta = \begin{bmatrix} \beta_0 \\ \beta_1 \\ \vdots \\ \beta_n \end{bmatrix}, \quad \mathbf{y} = \begin{bmatrix} y_0 \\ y_1 \\ \vdots \\ y_m \end{bmatrix}$$

and in matrix notation:

$$\mathbf{X}\beta = \mathbf{y}$$

A solver for system of linear equations based on [linear algebra](#) is the “*Normal equation*”.

This is the closed-form solution for the matrix equation  $\mathbf{X}\beta = \mathbf{y}$ :

$$\beta = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

where  $(\mathbf{X}^T \mathbf{X})$  is known as the “*Gram*” or “*normal*” matrix.

In python code:

```
X1 = np.column_stack(( X, np.ones(len(X)) ))
betas = np.linalg.inv(X1.T @ X1) @ X1.T @ y
print("beta_n,...,beta_0", betas)
```

## Polynomial regression

Not always (...basically never!) is the relationship between  $x$  and  $y$  is a straight line. We can add power terms *i.e.*

$$\hat{y}(\mathbf{X}) = \beta_2 \mathbf{X}^2 + \beta_1 \mathbf{X} + \beta_0$$

for example via

```
df["x_squared"] = df["x"]**2
```


and with this new feature we introduce curvature into our model

Note: This is still a linear function in terms of  $\beta$ , i.e. the fruits

			total
9	3	3	21€
25	5	2	20€
64	8	6	46€

# Extrapolation

**Warning!** Regression with high degree polynomials have unstable ‘tails’:



(Interpolation: within the **convex hull**, and extrapolation: outside)

## Linear regression using the scikit-learn

```
from sklearn.linear_model import LinearRegression
regressor = LinearRegression(fit_intercept=True)
# fit the model
regressor.fit(X, y)
# make predictions
y_pred = regressor.predict(X_test)
```

(The `fit_intercept=True` option will create the column of 1's needed to calculate  $\beta_0$ )

To tablulate the  $\beta$ 's

```
import eli5
feature_names = X.columns.tolist()
eli5.show_weights(regressor,
                  top=None,
                  feature_names = feature_names )
```

In Colab you will need to: `!pip install -q eli5`

and using the `statsmodels` library

(In Colab you will need to: `!pip install -q statsmodels`)

```
import statsmodels.api as sm

X = sm.add_constant(X) # add the column of 1's
model = sm.OLS(y,X)
results = model.fit()
print(results.summary())
```

# Explainability


The  $\beta_1$ , i.e. the inclinations, provide valuable information regarding the relative importance (weight) of each feature.

For example, if a feature has  $\beta_1 = 0$ , then this feature does not form part of the model, and has no influence when it comes to predicting  $\hat{y}$ .

Note: For a correct interpretation of the relative importance of the regression coefficients ( $\beta$ ) it is important that all of the features are on the same scale.

One can calculate the `confidence intervals (CI)` of the coefficients *i.e.* by using `statsmodels.regression.linear_model.OLS` with `conf_int`

To assist with our visual inspection ('EDA') we can use the [seaborn.lmplot](#)




which can provide the confidence interval ([ci](#)) of the fit

```
sns.lmplot(x=X, y=y, data=data, ci=95)
```

## The four basic assumptions of a good linear regression:


- Suffice to say, linear regression works best when the relation between the independent variable ( $x$ ) and the dependant variable( $y$ ) is actually linear



furthermore, to obtain reliable results for the regression coefficients and the confidence intervals the residuals (errors) should be:

- **i.i.d**: each error is independent of the other errors  
(no autocorrelation)
- **belong to a Gaussian distribution** ( $\forall i \in n, \varepsilon_i \sim \mathcal{N}(0, \sigma^2)$ )
- **homoscedasticity**: the errors have a constant variance ( $\sigma^2 \neq f(x)$ )

## Example of heteroscedasticity



Possible ‘solution’: apply a transformation, for example take the log or square root of the dependent variable ( $y$ )

(Note: this is not such a good idea if there is more than one independent variable)

# Loss and cost


How does the algorithm find the optimal solution?:

The *loss* ( $\mathcal{L}$ ) and *cost* ( $J$ ) functions

- *loss* is a distance measure\* between the points  $\hat{y}$  and  $y$
- *cost* is the mean of all of the losses:  $\bar{\mathcal{L}}$

(sometimes the loss function is known as the *error function*, in other words how bad are we doing  
and the cost functions is known as the *objective function*)

\*(Advanced material: norms: *p*-norm and Lebesgue spaces)



A valid loss function should have the following properties; it should be symmetric, *i.e.*

$$L(\hat{y}, y) = L(y, \hat{y})$$

it should be positive


$$L(\hat{y}, y) > 0, \quad \hat{y} \neq y$$

and should satisfy the triangle inequality

$$L(\hat{y}, y) \leq L(\hat{y}, z) + L(z, y)$$

*i.e.*  $L(\hat{y}, y)$  is the shortest distance between  $\hat{y}$  and  $y$ .

# What do the loss functions look like?



- L1 loss is not very sensitive to ‘outliers’
- L2 loss is sensitive to outliers as the distance is squared

In code:

$$y \rightarrow \text{y_true}$$

$$\hat{y} \rightarrow \text{y_pred}$$

$(\text{y_true} - \text{y_pred})$  is known as the error ( $\varepsilon$ ) or “*residual*”.

- L1 loss = AE (absolute error) =  $\text{abs}(\text{y_true} - \text{y_pred})$
- L1 cost = MAE (mean absolute error) =  $\frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$
- L2 loss = SE (squared error) =  $(\text{y_true} - \text{y_pred})^{**2}$
- L2 cost = MSE (mean squared error) =  $\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$

$$\text{RMSE} (\text{root mean squared error}) = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

We wish to minimize the cost (or ‘*objective function*’)  $J$

For example, in linear regression we use ordinary least squares, in other words the L2 loss and the L2 cost:

$$J(\beta) = MSE$$


$$= \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

$$J(\beta_1, \beta_0) = \frac{1}{n} \sum_{i=1}^n (y_i - \beta_1 \mathbf{X}_i + \beta_0)^2$$

The optimal values of  $\beta_1$  and  $\beta_0$  are given by  $\text{argmin}(J)$

(If we use the L1 loss we perform *median regression*, see [quantile regression](#))


What form does the cost ( $J$ ) have for the L1 loss for two points?



- $J$  is piecewise-continuous
- the minima is multivalued

(The mean absolute error is a new member of the collection of objectives in XGBoost (October 2022))


What form does the cost ( $J$ ) have for the L2 loss for two points?



- $J$  is smooth
- it is a **convex function** → has a unique global minima.

What form does the cost ( $J$ ) have for the L2 loss in 2d?

$J$  is now a parabolic surface with the same dimension as the number of parameters:



# Gradient descent

Solver: **gradient descent** (Augustin Louis Cauchy, 1847)

The gradient ( $G$ ) of a surface is given by

$$G = \frac{\partial J(\beta)}{\partial \beta}$$

where  $\beta$  are the parameters

Gradient descent '*update rule*'

$$\beta \leftarrow \beta - \eta \frac{\partial J(\beta)}{\partial \beta}$$

where  $\eta$  is the "*learning rate*" parameter

We wish to locate the local minima, or better still the global minima


(Paper: "*An overview of gradient descent optimization algorithms*")

algorithm:

first choose a random starting point for  $\beta$

- ① calculate the gradient ( $\partial J / \partial \beta$ )
- ② multiply by  $-1$  to change the sign, hence direction
- ③ scale by  $\eta$
- ④ this is the new value of  $\beta$
- ⑤ repeat, and when  $\beta$  hardly changes, stop

## Effect of the learning rate:




One can find a good  $\eta$  by plotting the cost step by step.


*Early stopping:* for example if  $\Delta J < 0,001$

If the cost rises then there is a big problem, and one must use a smaller  $\eta$


$\eta$  too small



good  $\eta$



$\eta$  too big



- Adagrad - Adaptive Gradient Algorithm: changes  $\eta$  on the fly

Gradient descent for linear regression:

$$\beta \leftarrow \beta - \eta \frac{\partial J(\beta)}{\partial \beta}$$

where

$$\begin{aligned} J(\beta) &= \text{MSE} = \frac{1}{n} (\varepsilon^T \varepsilon) = \frac{1}{n} (\hat{\mathbf{y}} - \mathbf{y})^T (\hat{\mathbf{y}} - \mathbf{y}) \\ &= \frac{1}{n} (\beta \mathbf{X} - \mathbf{y})^T (\beta \mathbf{X} - \mathbf{y}) \end{aligned}$$

leading to

$$\frac{\partial J(\beta)}{\partial \beta} = \frac{1}{n} ((\beta \mathbf{X} - \mathbf{y}) \mathbf{X})$$

giving


$$\beta \leftarrow \beta - \frac{\eta}{n} (\varepsilon \mathbf{X})$$

# Practical session

Let us look at gradient descent in our notebook

`notebook_ML_regression_GD.ipynb`

What happens if the features have very different scales?



## Feature scaling or normalization

With parametric estimators the solvers have better numerical stability if the data is “*normalized*”


Here are some linear transformations:

- `sklearn.preprocessing.MinMaxScaler`  $[0, 1]$
- `sklearn.preprocessing.StandardScaler`  $\rightarrow \mu = 0, \sigma = 1$
- `sklearn.preprocessing.RobustScaler` in case there are outliers
- (for neural networks there is the new `SquashingScaler` available in the `skrub` library)

Usage example:


```
from sklearn.preprocessing import RobustScaler  
  
RS = RobustScaler()  
  
X_train = RS.fit_transform(X_train)  
X_test = RS.transform(X_test)
```

# An example of a real neural network loss surface



(Source: "[Visualizing the Loss Landscape of Neural Nets](#)")

Linear regression in two dimensions, i.e. having two features:



The line is now a plane.

# Further reading

Free books on linear regression

- “The Truth about Linear Regression” by Cosma Rohilla Shalizi
- “Lecture notes on Ridge regression” by Wessel N. van Wieringen
- “Linear Model and Extensions” by Peng Ding
- “Regression and Other Stories” by Gelman, Hill, and Vehtari